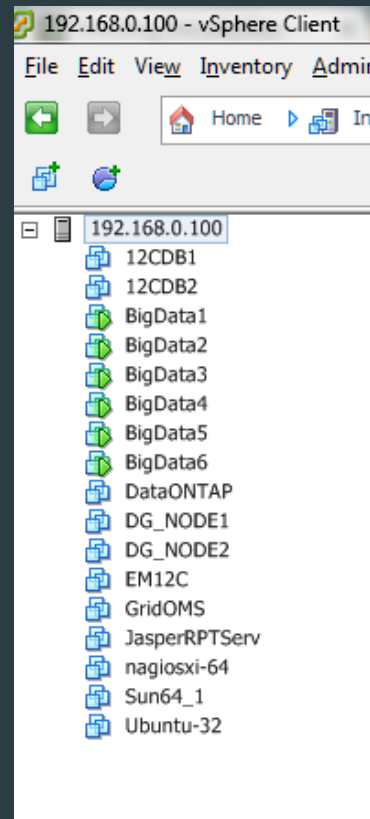


Configuring and Deploying MongoDB Sharded Cluster in 30 minutes

Prepared by Sudheer Kondla
Solutions Architect

- ▶ The presentation is based on VMs created using VMWare's ESXi/Vsphere and RHEL/Oracle Linux.
- ▶ This presentation is based on setting up mongoDB sharded cluster
 - ▶ with 6 VMs in the MongoDB cluster (RedHat Linux).
 - ▶ Each VM consists of 2 vCPUs and 4 GB of RAM
 - ▶ Each VM is created with 80 GB of disk space. No special mounts/ file systems are used.
 - ▶ Linux version used: 6.5

Sharded Cluster VIRTUAL MACHINES



VM Configuration

BigData1 - Virtual Machine Properties

Hardware | Options | Resources | Virtual Machine Version: 8

☐ Show All Devices Add... Remove

Hardware	Summary
Memory	4096 MB
CPUs	2
Video card	Video card
VMCI device	Restricted
SCSI controller 0	Paravirtual
Hard disk 1	Virtual Disk
CD/DVD drive 1	[VMSTORE1] V41362-0...
Network adapter 1	VM Network
Network adapter 2	VM Network
Floppy drive 1	Client Device

Memory Configuration

Memory Size: 4 GB

1011 GB
512 GB
256 GB
128 GB
64 GB
32 GB
16 GB
8 GB
4 GB
2 GB
1 GB

- Maximum recommended for this guest OS: 1011 GB.
- Maximum recommended for best performance: 24568 MB.
- Default recommended for this guest OS: 2 GB.
- Minimum recommended for this guest OS: 512 MB.

Installing mongodb

- ▶ The following steps guide you through installing mongodb software on Linux.
- ▶ set yum repository and download packages using yum package installer.
- ▶ `[root@bigdata1 ~]# cd /etc/yum.repos.d/`
 - ▶ `[root@bigdata1 yum.repos.d]# cat mongodb.repo`

```
[mongodb]
name=MongoDB Repository
baseurl=http://downloads-distro.mongodb.org/repo/redhat/os/x86_64/
gpgcheck=0
enabled=1
```
- ▶ Before you run “yum install”, be sure to check internet is working.
- ▶ `[root@bigdata6 yum.repos.d]# yum install -y mongodb-org-2.6.1 mongodb-org-server-2.6.1 mongodb-org-shell-2.6.1 mongodb-org-mongos-2.6.1 mongodb-org-tools-2.6.1`

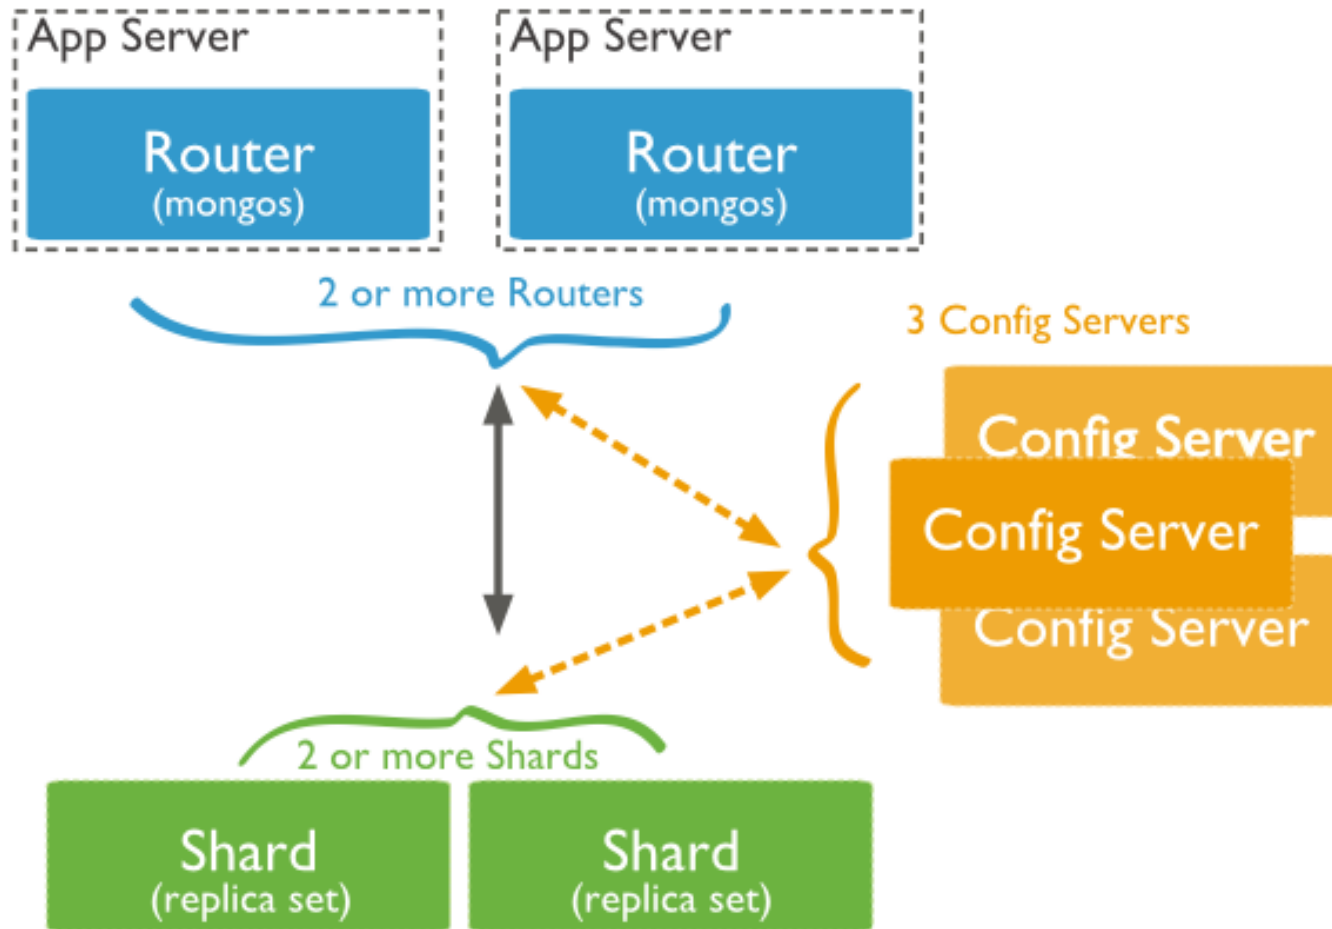
Setting up mongod

- ▶ Make sure to create /data/configdb and /data/db directories on each servers
- ▶ The above directories should be owned by mongod user and mongod group.
- ▶ Change ownership to mongod
- ▶ As a root user run “chown mongod:mongod /data/configdb” and “chown mongod:mongod /data/db”
- ▶ Without above directories mongod process will not start
- ▶ When you start mongo daemon process it will create “/data/configdb/mongod.lock” file
- ▶ You can also start mongod process with service option as root. For example “service mongod start”
- ▶ You can also configure mongo daemon to start at system boot.

MongoDB Sharding Architecture

- ▶ Sharding is a method for storing data across multiple machines. MongoDB uses sharding to support deployments with very large data sets and high throughput operations.
- ▶ Database systems with large data sets and high throughput applications can challenge the capacity of a single server. High query rates can exhaust the CPU capacity of the server. Larger data sets exceed the storage capacity of a single machine. Finally, working set sizes larger than the system's RAM stress the I/O capacity of disk drives.
- ▶ Sharding addresses the challenge of scaling to support high throughput and large data sets:
 - ▶ Sharding reduces the number of operations each shard handles. Each shard processes fewer operations as the cluster grows. As a result, a cluster can increase capacity and throughput horizontally.
 - ▶ For example, to insert data, the application only needs to access the shard responsible for that record.
 - ▶ Sharding reduces the amount of data that each server needs to store. Each shard stores less data as the cluster grows.
 - ▶ For example, if a database has a 1 terabyte data set, and there are 4 shards, then each shard might hold only 256GB of data. If there are 40 shards, then each shard might hold only 25GB of data.

Sharding Architecture



Deploying a sharded cluster

- ▶ Start the Config Server Database Instances.
 - ▶ The config server processes are mongod instances that store the cluster's metadata.
 - ▶ Designate a mongod as a config server using the `--configsvr` option.
e.g: `mongod --configsvr --dbpath /data/configdb --port 27019 --bind_ip 192.168.0.131 -v`
 - ▶ Each config server stores a complete copy of the cluster's metadata.
 - ▶ Production deployments require three config server instances, each of them running on different servers to ensure high availability and data safety.
 - ▶ All members of a sharded cluster must be able to connect to all other members of a sharded cluster.
 - ▶ Make sure to setup proper network connectivity, firewall rules.
 - ▶ Start the three config server instances. For example
 - ▶ Node1: `mongod --configsvr --dbpath /data/configdb --port 27019 --bind_ip 192.168.0.131 -v`
 - ▶ Node2: `mongod --configsvr --dbpath /data/configdb --port 27019 --bind_ip 192.168.0.132 -v`
 - ▶ Node3: `mongod --configsvr --dbpath /data/configdb --port 27019 --bind_ip 192.168.0.133 -v`
- ▶ Start the mongos Instances:
 - ▶ The mongos instances are lightweight and do not require data directories.
 - ▶ You can run a mongos instance on a system that runs other cluster components, such as on an application server or a server running a mongod process.
 - ▶ By default, a mongos instance runs on port 27017. Specify the hostnames of the three config servers, either in the configuration file or as command line parameters.
 - ▶ For example: `mongos --configdb bigdata1:27019,bigdata2:27019,bigdata3:27019`

Add Shards to the Cluster

► Enable Sharding for a Database

- Enabling sharding for a database does not redistribute data but make it possible to shard the collections in that database.
- From a mongo shell, connect to the mongos instance. Issue a command using the following

```
[hdfs@bigdata3 ~]$ mongo --host bigdata2 --port 27017
```

```
MongoDB shell version: 2.6.1
```

```
connecting to: bigdata2:27017/test
```

```
rs0:PRIMARY>
```

- `sh.addShard("rs0/bigdata1:27017")`
- `sh.addShard("rs0/bigdata2:27017")`
- `sh.addShard("rs0/bigdata3:27017")`
- `sh.addShard("rs0/bigdata4:27017")`
- `sh.addShard("rs0/bigdata5:27017")`
- `sh.addShard("rs0/bigdata6:27017")`

Enable Sharding for a Collection

- ▶ Before you can shard a collection, you must enable sharding for the collection's database.
- ▶ Enabling sharding for a database does not redistribute data but make it possible to shard the collections in that database.
- ▶ Once you enable sharding for a database, MongoDB assigns a primary shard for that database where MongoDB stores all data before sharding begins.
- ▶ Issue the `sh.enableSharding()` method.
 - ▶ `rs0:PRIMARY> show dbs`
 - ▶ `sh.enableSharding("customer")` OR
 - ▶ `db.runCommand({ enableSharding: "customer" })`
- ▶ Enable sharding on a per-collection basis.
 - ▶ `sh.shardCollection("records.people", { "zipcode": 1, "name": 1 })`
 - ▶ `sh.shardCollection("people.addresses", { "state": 1, "_id": 1 })`
 - ▶ `sh.shardCollection("assets.chairs", { "type": 1, "_id": 1 })`
 - ▶ `sh.shardCollection("events.alerts", { "_id": "hashed" })`